

Operating System

Chapter 6. Concurrency: Deadlock and Starvation



Lynn Choi

School of Electrical Engineering



高麗大學校

Computer System Laboratory



□ Definition

- A set of processes is deadlocked when each process in the set is blocked awaiting an event (or a resource) that can only be triggered (released) by another blocked process in the set


□ Examples

- 4 cars arrive at a four-way stop

□ Two general categories of resources


- Reusable resource
 - Can be safely used by only one process at a time and is not depleted by that use
 - Examples: processors, memory, I/O devices, files, databases and semaphores
- Consumable resource
 - Can be created (produced) and destroyed (consumed)
 - Typically, there is no limit on the number of consumable resources
 - Examples: interrupts, signals, messages, data in I/O buffers

Potential Deadlock



Source: Pearson

Actual Deadlock



Source: Pearson

Reusable Resource Example



Process P		Process Q	
Step	Action	Step	Action
p ₀	Request (D)	q ₀	Request (T)
p ₁	Lock (D)	q ₁	Lock (T)
p ₂	Request (T)	q ₂	Request (D)
p ₃	Lock (T)	q ₃	Lock (D)
p ₄	Perform function	q ₄	Perform function
p ₅	Unlock (D)	q ₅	Unlock (T)
p ₆	Unlock (T)	q ₆	Unlock (D)

Figure 6.4 Example of Two Processes Competing for Reusable Resources


p₀p₁q₀q₁p₂q₂ leads to a deadlock!

Source: Pearson

Consumable Resource Example



- ❑ Consider a pair of processes, in which each process attempts to receive a message from the other process and then send a message to the other process
- ❑ Deadlock occurs if the Receive is blocking



Source: Pearson

Conditions for Deadlock



❑ Mutual exclusion

- Only one process may use a resource at a time

❑ Hold and wait

- A process may hold allocated resources while waiting for the other resources

❑ No preemption

- No resource can be forcibly removed from a process holding it

❑ Circular wait

- A closed chain of processes exists such that each process holds at least one resource needed by the next process in the chain

❑ Note that

- The first three conditions are necessary but not sufficient conditions for a deadlock.
 - The fourth condition is actually a consequence of the first three.
- Given that the first three conditions exist, a sequence of events may occur that lead to an unresolvable circular wait.


Resource Allocation Graphs




(a) Resource is requested



(b) Resource is held



(c) Circular wait



(d) No deadlock

Source: Pearson

Circular Wait Example




Figure 6.6 Resource Allocation Graph for Figure 6.1b

Source: Pearson

Three Approaches for Deadlocks



❑ **Deadlock prevention**

- Adopt a policy that eliminates one of the conditions 1 through 4

❑ **Deadlock avoidance**

- Make the appropriate choices dynamically based on the current state of resource allocation

❑ **Deadlock detection**

- Allow the deadlock to occur, attempt to detect the presence of deadlock, and recover if a deadlock is detected

Deadlock Detection, Prevention, Avoidance



Approach	Resource Allocation Policy	Different Schemes	Major Advantages	Major Disadvantages
Prevention	Conservative; undercommits resources	Requesting all resources at once	<ul style="list-style-type: none"> • Works well for processes that perform a single burst of activity • No preemption necessary 	<ul style="list-style-type: none"> • Inefficient • Delays process initiation • Future resource requirements must be known by processes
		Preemption	<ul style="list-style-type: none"> • Convenient when applied to resources whose state can be saved and restored easily 	<ul style="list-style-type: none"> • Preempts more often than necessary
		Resource ordering	<ul style="list-style-type: none"> • Feasible to enforce via compile-time checks • Needs no run-time computation since problem is solved in system design 	<ul style="list-style-type: none"> • Disallows incremental resource requests
Avoidance	Midway between that of detection and prevention	Manipulate to find at least one safe path	<ul style="list-style-type: none"> • No preemption necessary 	<ul style="list-style-type: none"> • Future resource requirements must be known by OS • Processes can be blocked for long periods
Detection	Very liberal; requested resources are granted where possible	Invoke periodically to test for deadlock	<ul style="list-style-type: none"> • Never delays process initiation • Facilitates online handling 	<ul style="list-style-type: none"> • Inherent preemption losses

Source: Pearson

Deadlock Prevention



❑ Mutual exclusion

- We cannot prevent this first condition
 - If access to a resource requires mutual exclusion, then mutual exclusion must be supported by the system

❑ Hold and wait

- Require that a process request all of its required resources at once and block the process until all the requests can be granted simultaneously

❑ No preemption

- If a process holding certain resources is denied a further request, that process must release its original resources and request them again
- Alternatively, if a process requests a resource that is currently held by another process, OS may preempt the second process

❑ Circular wait

- Define a linear ordering of resource types
- If a process has been allocated resources of type R, then it may subsequently request only those resources of types following R in the ordering

Deadlock Avoidance



❑ Deadlock avoidance

- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
 - Requires knowledge of future resource requests

❑ Two approaches

- Process initiation denial
 - Do not start a process if its demands may lead to a deadlock
- Resource allocation denial
 - Do not grant a resource request to a process if this allocation might lead to a deadlock

❑ Advantages

- It is less restrictive than deadlock prevention
- It is not necessary to preempt and rollback processes, as in deadlock detection

❑ Require

- The maximum resource requirement must be known in advance

Process Initiation Denial



- Consider a system of n processes and m different types of resources
- Let's define the following vectors and matrices:

➤ Resource = $\mathbf{R} = (R_1, R_2, \dots, R_m)$

➤ Available = $\mathbf{V} = (V_1, V_2, \dots, V_m)$

➤ Claim = $\mathbf{C} = \begin{bmatrix} C_{11} & C_{12} & \dots & C_{1m} \\ C_{21} & C_{22} & \dots & C_{2m} \\ \cdot & \cdot & & \cdot \\ C_{n1} & C_{n2} & \dots & C_{nm} \end{bmatrix}$

➤ Allocation = $\mathbf{A} = \begin{bmatrix} A_{11} & A_{12} & \dots & A_{1m} \\ A_{21} & A_{22} & \dots & A_{2m} \\ \cdot & \cdot & & \cdot \\ A_{n1} & A_{n2} & \dots & A_{nm} \end{bmatrix}$

Process Initiation Denial



□ The following relationship holds

- $R_j = V_j + \sum_{i=1}^n A_{ij}$ for all j
 - All resources are either available or allocated.
- $C_{ij} \leq R_j$ for all i, j
 - No process can claim more than total amount of resources
- $A_{ij} \leq C_{ij}$ for all i, j
 - No process is allocated more resources than it originally claimed

□ Policy: start a new process P_{n+1} only if

- $R_j \geq C_{(n+1)j} + \sum_{i=1}^n C_{ij}$ for all j
 - A process is only started if the maximum claim of all current processes plus those of the new process can be met

Resource Allocation Denial



□ Referred to as the *banker's algorithm*

- State of the system reflects the current allocation of resources to processes
- Safe state is one in which there is at least one sequence of resource allocations to processes that does not result in a deadlock
- Unsafe state is a state that is not safe

Determination of a Safe State



- System state consists of 4 processes and 3 resources

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
0	1	1

Available vector V

(a) Initial state

Source: Pearson

- Is this a safe state?

- Can any of 4 processes run to completion?
 - P2 can run to completion!

P2 Runs to Completion



- After P2 completes, P2 releases its resources

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
6	2	3

Available vector V

(b) P2 runs to completion

Source: Pearson

- Then, we can run any of P1, P3, or P4
- Assume we select P1

P1 Runs to Completion



	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
7	2	3

Available vector V

(c) P1 runs to completion

Source: Pearson

P3 Runs to Completion



	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
9	3	4

Available vector V

(d) P3 runs to completion

Source: Pearson

Determination of an Unsafe State



	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	1	0	2
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
1	1	2

Available vector V

(a) Initial state

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	2	0	1
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	1	2	1
P2	1	0	2
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector R

R1	R2	R3
0	1	1

Available vector V

(b) P1 requests one unit each of R1 and R3

❑ Is this a safe state?

Source: Pearson

Deadlock Avoidance Logic



```
struct state {  
    int resource[m];  
    int available[m];  
    int claim[n][m];  
    int alloc[n][m];  
}
```

(a) global data structures

```
if (alloc [i,*] + request [*] > claim [i,*])  
    < error >; /* total request > claim*/  
else if (request [*] > available [*])  
    < suspend process >;  
else { /* simulate alloc */  
    < define newstate by:  
        alloc [i,*] = alloc [i,*] + request [*];  
        available [*] = available [*] - request [*] >;  
    }  
    if (safe (newstate))  
        < carry out allocation >;  
    else {  
        < restore original state >;  
        < suspend process >;  
    }  
}
```

(b) resource alloc algorithm

Source: Pearson

Deadlock Avoidance Logic



```
boolean safe (state S) {
    int currentavail[m];
    process rest[<number of processes>];
    currentavail = available;
    rest = {all processes};
    possible = true;
    while (possible) {
        <find a process  $P_k$  in rest such that
            claim  $[k,*] - \text{alloc } [k,*] \leq \text{currentavail};$ 
        if (found) {
            /* simulate execution of  $P_k$  */
            currentavail = currentavail + alloc  $[k,*]$ ;
            rest = rest -  $\{P_k\}$ ;
        }
        else possible = false;
    }
    return (rest == null);
}
```

(c) test for safety algorithm (banker's algorithm)

Figure 6.9 Deadlock Avoidance Logic

Source: Pearson

Deadlock Detection



- ❑ **Deadlock prevention is very conservative**
 - Limit access to resources by imposing restrictions on processes
- ❑ **Deadlock detection does the opposite**
 - Resource requests are granted whenever possible
- ❑ **A check for deadlock can be made**
 - As frequently as each resource request
 - (+) Lead to early detection
 - (+) Can employ a relatively simple detection algorithm because it is based on incremental changes to the system state
 - (-) Consume considerable processor time
 - Or less frequently depending on how likely it is for a deadlock to occur

Deadlock Detection Algorithm



- ❑ Instead of Claim (C), a Request (Q) matrix is defined
 - Q_{ij} represents the amount of resources of type j requested by process i
- ❑ Initially, all processes are unmarked (deadlocked)
- ❑ The algorithm proceeds by marking processes that are not deadlocked.
- ❑ Then, the following steps are performed
 1. Mark each process that has a row of all zeros in the Allocation matrix
 2. Initialize a temporary vector **W** to equal the Available vector
 3. Find an index i such that process is currently unmarked and the i^{th} row of Q is less than or equal to **W**. If no such row is found, terminate the algorithm
 4. If such a row is found, mark process i and add the corresponding row of the allocation matrix to **W**. Return to step 3
- ❑ A deadlock exists if and only if there are unmarked processes at the end of the algorithm

Deadlock Detection Algorithm



	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource vector

R1	R2	R3	R4	R5
0	0	0	0	1

Available vector

Figure 6.10 Example for Deadlock Detection

Source: Pearson

- Mark P4 because P4 has no allocated resources
- Set $\mathbf{W} = (0 \ 0 \ 0 \ 0 \ 1)$
- The request of P3 is less than or equal to W, so mark P3 and set
 - $\mathbf{W} = \mathbf{W} + (0 \ 0 \ 0 \ 1 \ 0) = (0 \ 0 \ 0 \ 1 \ 1)$
- No other unmarked process has a row in Q that is less than or equal to W. Therefore, terminate the algorithm. P1 and P2 are deadlocked!

Deadlock Recovery



□ Recovery options in order of increasing sophistication

- Abort all deadlocked processes.
 - The most common solution adopted by OS
- Backup each deadlocked process to a previous checkpoint and restart
 - Require rollback and restart mechanism
 - The original deadlock may recur
 - However, the nondeterminism of concurrent processing may ensure that this does not happen
- Successively abort deadlocked process until deadlock no longer exists
 - After each abortion, the detection algorithm must be re-invoked
- Successively preempt resources until deadlock no longer exists
 - A process that has a resource preempted must be rolled back to a prior point before its acquisition of the resource

Dining Philosophers Problem



- ❑ No two philosophers can use the same fork at the same time
 - Mutual exclusion
- ❑ No philosopher must starve to death
 - Avoid starvation and deadlock




Figure 6.11 Dining Arrangement for Philosophers

Source: Pearson

Solution I using Semaphore



```
/* program      diningphilosophers */
semaphore fork [5] = {1};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal(fork [(i+1) mod 5]);
        signal(fork[i]);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher
(2),
            philosopher (3), philosopher (4));
}
```

Figure 6.12 A First Solution to the Dining Philosophers Problem

Source: Pearson

Solution II using Semaphore



```
/* program diningphilosophers */
semaphore fork[5] = {1};
semaphore room = {4};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (room);
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal (fork [(i+1) mod 5]);
        signal (fork[i]);
        signal (room);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1), philosopher (2),
              philosopher (3), philosopher (4));
}
```

Figure 6.13 A Second Solution to the Dining Philosophers Problem

Source: Pearson

Solution using Monitor

```
monitor dining_controller;
cond ForkReady[5];          /* condition variable for synchronization */
boolean fork[5] = {true};    /* availability status of each fork */

void get_forks(int pid)      /* pid is the philosopher id number */
{
    int left = pid;
    int right = (++pid) % 5;
    /*grant the left fork*/
    if (!fork(left))
        cwait(ForkReady[left]);          /* queue on condition variable */
    fork(left) = false;
    /*grant the right fork*/
    if (!fork(right))
        cwait(ForkReady[right]);         /* queue on condition variable */
    fork(right) = false;
}
void release_forks(int pid)
{
    int left = pid;
    int right = (++pid) % 5;
    /*release the left fork*/
    if (empty(ForkReady[left])          /*no one is waiting for this fork */
        fork(left) = true;
    else
        csignal(ForkReady[left]);       /* awaken a process waiting on this fork */
    /*release the right fork*/
    if (empty(ForkReady[right])          /*no one is waiting for this fork */
        fork(right) = true;
    else
        csignal(ForkReady[right]);      /* awaken a process waiting on this fork */
}

void philosopher[k=0 to 4]      /* the five philosopher clients */
{
    while (true) {
        <think>;
        get_forks(k);             /* client requests two forks via monitor */
        <eat spaghetti>;
        release_forks(k);         /* client releases forks via the monitor */
    }
}
```

Source: Pearson

Figure 6.14 A Solution to the Dining Philosophers Problem Using a Monitor

UNIX Concurrency Mechanisms



- ❑ UNIX provides a variety of mechanisms for interprocess communication and synchronization including:
 - **Pipes**
 - First-in-first-out queue, written by one process and read by another
 - Implemented by a circular buffer, allowing two processes to communicate on the producer-consumer model
 - Example: `ls | more`, `ps | sort`, etc.
 - **Messages**
 - UNIX provides *msgsnd* and *msgrcv* system calls for processes to engage in message passing
 - A message is a block of bytes
 - Each process is associated a *message queue*, which functions like a mailbox

UNIX Concurrency Mechanisms



➤ Shared memory

- Common block of virtual memory shared by multiple processes
- Fastest form of interprocess communication
- Mutual exclusion is provided for each location in shared-memory
- A process may have read-only or read-write permission for a memory location

➤ Semaphores

- Generalization of the semWait and semSignal primitives defined in Chapter 5
- Increment and decrement operations can be greater than 1
 - ▼ Thus, a single semaphore operation may involve incrementing/decrementing a semaphore and waking up/suspending processes.
 - ▼ Provide considerable flexibility in process synchronization

UNIX Concurrency Mechanisms



➤ Signals

- A software mechanism that informs a process of the occurrence of asynchronous events (similar to a hardware interrupt)
- Sending a signal
 - ▼ Kernel *sends* (delivers) a signal to a *destination process* by updating some state in the context of the destination process.
 - ▼ Kernel sends a signal for one of the following reasons:
 - ◆ Kernel has detected a system event such as divide-by-zero (SIGFPE) or the termination of a child process (SIGCHLD)
 - ◆ Another process has invoked the kill system call to explicitly request the kernel to send a signal to the destination process.
- Receiving a signal
 - ▼ A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal.
 - ▼ Two possible ways to react:
 - ◆ Default action (ignore, terminate the process, terminate & dump)
 - ◆ *Catch* the signal by executing a user-level function called a *signal handler*.
 - ▼ Akin to a hardware exception handler being called in response to an asynchronous interrupt.

UNIX Signals



Value	Name	Description
01	SIGHUP	Hang up; sent to process when kernel assumes that the user of that process is doing no useful work
02	SIGINT	Interrupt
03	SIGQUIT	Quit; sent by user to induce halting of process and production of core dump
04	SIGILL	Illegal instruction
05	SIGTRAP	Trace trap; triggers the execution of code for process tracing
06	SIGIOT	IOT instruction
07	SIGEMT	EMT instruction
08	SIGFPE	Floating-point exception
09	SIGKILL	Kill; terminate process
10	SIGBUS	Bus error
11	SIGSEGV	Segmentation violation; process attempts to access location outside its virtual address space
12	SIGSYS	Bad argument to system call
13	SIGPIPE	Write on a pipe that has no readers attached to it
14	SIGALRM	Alarm clock; issued when a process wishes to receive a signal after a period of time
15	SIGTERM	Software termination
16	SIGUSR1	User-defined signal 1
17	SIGUSR2	User-defined signal 2
18	SIGCHLD	Death of a child
19	SIGPWR	Power failure

Source: Pearson

Signal Concepts (cont)



- ❑ A signal is **pending** if it has been sent but not yet received.
 - There can be at most one pending signal of any particular type.
 - Important: Signals are not queued
 - If a process has a pending signal of type *k*, then subsequent signals of type *k* that are sent to that process are discarded.
- ❑ A process can **block** the receipt of certain signals.
 - Blocked signals can be delivered, but will not be received until the signal is unblocked.
- ❑ Kernel maintains **pending** and **blocked** bit vectors in the context of each process.
 - `pending` – represents the set of pending signals
 - Kernel sets bit *k* in `pending` whenever a signal of type *k* is delivered.
 - Kernel clears bit *k* in `pending` whenever a signal of type *k* is received
 - `blocked` – represents the set of blocked signals
 - Can be set and cleared by the application using the `sigprocmask` function.

Receiving Signals



- ❑ Suppose kernel is returning from exception handler and is ready to pass control to process p .
- ❑ Kernel computes $\text{pnb} = \text{pending} \ \& \ \sim\text{blocked}$
 - The set of pending nonblocked signals for process p
- ❑ If $(\text{pnb} == 0)$
 - Pass control to next instruction in the logical flow for p .
- ❑ Else
 - Choose least nonzero bit k in pnb and force process p to **receive** signal k .
 - The receipt of the signal triggers some **action** by p
 - Repeat for all nonzero k in pnb .
 - Pass control to next instruction in logical flow for p .

Installing Signal Handlers



- ❑ The `signal` function modifies the default action associated with the receipt of signal `signum`:

- `handler_t *signal(int signum, handler_t *handler)`

- ❑ Different values for `handler`:

- `SIG_IGN`: ignore signals of type `signum`
 - `SIG_DFL`: revert to the default action on receipt of signals of type `signum`.
 - Otherwise, `handler` is the address of a user-defined function called *signal handler*
 - Called when process receives signal of type `signum`
 - Changing the default action by passing the address of a handler to the `signal` function is known as “*installing*” the handler.
 - The invocation of the handler is called “*catching*” the signal
 - The execution of the handler is referred as “*handling*” the signal.
 - When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal.

Signal Handling Example



```
#include "csapp.h"

void handler(int sig) {
    static int beeps = 0;
    printf("BEEP\n");
    if (++beeps < 5) Alarm(1); /* next SIGALRM
will be delivered in 1s */
    else {
        printf("BOOM!\n");
        exit(0);
    }
}

int main() {

    Signal(SIGALRM, handler); /* install
SIGALRM handler */
    Alarm(1); /* next SIGALRM will be delivered
in 1s */
    while (1) { ; /* signal handler returns
control here each time */ }
    exit(0);
}
```

```
linux> ./alarm
BEEP
BEEP
BEEP
BEEP
BEEP
BOOM!
linux>
```

Source: Pearson

Homework 5



- ☐ Exercise 6.5
- ☐ Exercise 6.8
- ☐ Exercise 6.9
- ☐ Exercise 6.11
- ☐ Exercise 6.14